

Question 1: For a specific RAID array (call it “RAID A”), a read of a block takes about 10ms. A write of a block also takes about 10ms. This RAID is likely:

- a) RAID-1 (mirroring)
- b) RAID-4 (parity disk)
- c) RAID-5 (rotating parity) d) RAID-4 or RAID-5
- e) All of the above

Question 2: for “RAID B”, two small random writes usually take about twice as long as one random write. This RAID is likely:

- a) RAID-1 (mirroring)
- b) RAID-4 (parity disk)
- c) RAID-5 (rotating parity) d) RAID-4 or RAID-5
- e) All of the above

Question 3: for “RAID C”, a large write (of 7 blocks) usually takes about as much time as a small write (1 block). This RAID is likely:

- a) RAID-1 (mirroring)
- b) RAID-4 (parity disk)
- c) RAID-5 (rotating parity) d) RAID-4 or RAID-5
- e) All of the above

Question 4: For “RAID D”, the overall throughput (measured in MB/s) is about 4 MB/s when issuing many 1- block random writes. In comparison, a comparable RAID array configured to use striping (RAID-0) achieved a throughput of about 8 MB/s. This RAID (RAID D) is likely:

- a) RAID-1 (mirroring)
- b) RAID-4 (parity disk)
- c) RAID-5 (rotating parity) d) RAID-4 or RAID-5
- e) All of the above

Question 5: for “RAID E”, the overall throughput when writing large sequential blocks to disk is about 700 MB/ s. In comparison, large sequential writes to a RAID 0 achieves about 800 MB/s. This RAID (RAID-E) is likely: a) RAID-1 (mirroring)

- b) RAID-4 (parity disk)
- c) RAID-5 (rotating parity)
- d) RAID-4 or RAID-5
- e) All of the above

For the next part of the RAID question, you discover a new forensic tool: the ability to measure the number of physical I/Os that happen in the system (you do this by attaching probes to the internal I/O busses of the RAID array in question; nice work!).

Question 6: For “RAID F”, you know that the RAID is likely RAID-4 or RAID-5. You issue a single perfectly aligned block write. The number of physical I/Os you measure on RAID F during this write is:

- a) 0
- b) 1
- c) 2
- d) 3
- e) 4

Question 7: For “RAID G”, you have already figured out that it is likely a RAID-1 (mirroring). You then issue a single write to the RAID. The write is small (1 block) is aligned. The number of physical I/Os you measure on the disks of RAID G is always:

- a) 0
- b) 1
- c) 2
- d) Sometimes =2, sometimes >2
- e) Always > 2

Question 8: For “RAID H”, you have already figured out that it is likely a RAID-1 (mirroring). You then issue a single write to the RAID. The write is small (1 block); however, it is not necessarily aligned. The number of physical I/Os you measure on the disks of RAID H is always:

- a) 0
- b) 1
- c) 2
- d) Sometimes =2, sometimes >2
- e) Always > 2

Question 9: For “RAID I”, you already know that it is RAID-5 (rotating parity). You then issue a single small read (1 block), which is never aligned. The number of physical I/Os you measure on the disks of RAID I is always:

- a) 0
- b) 1
- c) 2
- d) Sometimes =2, sometimes >2
- e) Always > 2

Question 10: For “RAID J”, you measure the total number of physical I/Os under a read-only workload. You find that it is equal to the number of logical reads issued to the RAID. You also see that one disk is never accessed. From this, you conclude that the RAID is:

- a) RAID-0 (striping)
- b) RAID-1 (mirroring)
- c) RAID-4 (parity disk)
- d) RAID-0 or RAID-1
- e) RAID-0 or RAID-1 or RAID-4

A new type of disk has come to market. It is called a “shingled disk”, and has the following properties. First, tracks are grouped into groups called “shingles”. Within each shingle, tracks are packed very tightly, such that when you write to a block on track K, you will likely overwrite (oops) the corresponding block on track K+1. Thus, tracks within a shingle should be written in order, from 0 to T-1 (assuming T tracks in each shingle);

if you stick to writing tracks in order within a shingle, you won’t lose any data.

The reason for shingles: it allows disk manufacturers to pack even more data on disk, making disks cheaper.

The problem: you can’t easily overwrite a block in place; rather, the device needs to keep a mapping table of some kind (the Shingle Translation Layer, or STL), and write blocks out (within each shingle) in log-structured fashion.

Question 11: Assume the STL maps each 4KB block to a location on the disk. How big must the STL be to hold translations for an entire 512-GB disk? Assume a 4-byte disk address for each entry in an array-like structure in the STL.

- a) 1 MB
- b) 128 MB
- c) 512 MB
- d) 1 GB
- e) None of the above

Question 12: The STL size can be changed by mapping chunks in sizes other than the usual 4KB block. With each doubling of the block size, the STL:

- a) Increases in size by 2x
- b) Decreases in size by 2x
- c) Increases in size by 4x
- d) Decreases in size by 4x
- e) None of the above

Question 13: In a shingled disk, all writes are log-structured. As a result, which of the following is NOT true: a) Write performance is similar to a regular hard drive for sequential workloads

- b) Write performance is similar to a regular hard drive for random workloads
- c) Read performance is similar to a regular hard drive for sequential workloads
- d) Read performance is similar to a regular hard drive for random workloads
- e) All of the above are true

Question 14: The STL size can be changed by mapping chunks in sizes other than the usual 4KB block. With each doubling of the block size, the STL:

- a) Increases in size by 2x
- b) Decreases in size by 2x
- c) Increases in size by 4x
- d) Decreases in size by 4x
- e) None of the above

Question 15: In a shingled disk, all writes are log-structured. As a result, which one of these is NOT true: a) Performance is similar to a regular hard drive for sequential write workloads

- b) Performance is similar to a regular hard drive for random write workloads
- c) Performance is similar to a regular hard drive for sequential read workloads
- d) Performance is similar to a regular hard drive for random read workloads
- e) Performance is similar to a regular hard drive for workloads with a mix of sequential reads/writes

With this basic knowledge in place, let us now explore some forensics questions about a shingled drive you have.

Question 16: You have a suspicion that the STL cannot hold all mapping entries in device memory; as a result, sometimes when reading from disk, an extra read is incurred (to look up the mapping information). The workload used here is to repeatedly read some number of blocks N (many times), and report the average time per read; N is increased along the x-axis. Which line on the graph best reflects this result?

- a) Lowest line with drop halfway
- b) Next lowest dashed line with drop halfway
- c) Solid line with rise halfway
- d) Topmost line which goes up then down
- e) None of the above

Question 17: Assuming a FIFO scheduling policy, and a request stream which reads blocks 21, 17, 11, 7, 13, what is the last block to be read?

- a) 21
- b) 17
- c) 11
- d) 7
- e) 13

Question 18: Same question, but assume a SSTF scheduling policy, and a VERY FAST seek. What is the last block read?

- a) 21
- b) 17
- c) 11
- d) 7
- e) 13

Question 19: Same question, but assume SATF scheduling policy, and a VERY FAST seek. What is the last block read?

- a) 21
- b) 17
- c) 11
- d) 7
- e) 13

Question 20: In a shingled disk, all writes are log-structured. Specifically, writes are directed to the currently- being-written shingle, and the STL is updated accordingly. As a result, does disk scheduling of write requests (assuming a workload only contains write requests) help performance?

- a) Yes (always)
- b) Yes (sometimes)
- c) No (never)
- d) Can't answer without more details
- e) None of the above

Question 21: A file descriptor is:

- a) a system-wide object used to access files
- b) a per-process integer used to access files
- c) readily forged
- d) returned to a process via the close() system call
- e) hard to understand

Question 22: Adding the O_TRUNC flag to the open() call will

- a) puts the data in a locked "trunk" file
- b) causes the open call to fail (usually)
- c) implies you must add O_TRUNC to close() as well
- d) creates the file (if it doesn't exist)
- e) truncates the file to size=0

Question 23: The "unlink()" system call is to the program "rm" as THIS is to the "rmdir" program

- a) unlink()
- b) delete()
- c) rmdir()
- d) link()
- e) fork()

Question 24: The "lseek()" call is used to

- a) reposition the disk head
- b) do a long disk seek, immediately
- c) change the current file offset
- d) force changes to disk
- e) close files after a layoff period

Question 25: The following information is NOT available within a typical inode:

- a) owner
- b) size (bytes)
- c) blocks allocated
- d) file name
- e) last access time

Question 26: The “read()” call is to the “cat” program as the BLANK call is to the “ls” program:

- a) read()
- b) readdir()
- c) stat()
- d) fstat()
- e) umount()

Question 27: Which type of links do most UNIX file systems support?

- a) hyperlinks
- b) soft links
- c) forked links
- d) sausage links
- e) unlinks

Question 28: A file is NOT

- a) a container for data
- b) a byte array that can be read or written
- c) something with a low-level name
- d) easily deleted
- e) something that can be referred to via a high-level name, thanks to directories

Question 29: Let’s say we wish to write data to a file and then force the contents of a file to disk. We should thus call:

- a) write()
- b) write() then fsync()
- c) write() then fopen()
- d) write() then fallocc()
- e) write() then fit()

Question 30: To atomically replace the contents of a file foo, we should use the following sequence of system calls: open() [to open a new temporary file], write() to write data to the disk, force the contents to disk (with a certain file system call, as in question 29), and finally:

- a) close() to close the file
- b) rename() to change the name of the temporary file to the desired file name
- c) link() to link the file to another name
- d) unlink() to remove the original file foo
- e) None of the above

In this bit of forensics, only bits and pieces remain on disk of a file system that looks very similar to the Very Simple File System (VSFS). VSFS, as you may recall, has an inode bitmap, some inodes, a data bitmap, and some data blocks. Here is a sample VSFS empty file system disk image:

```
Inode bitmap 10000000
Inodes [d a:0 r:2] [] [] [] [] []
data bitmap 10000000
data [(.,0) (.,0)] [] [] [] [] []
```

As you can see, only the (empty) root directory is allocated in this tiny file system. In the inode of the root directory, you see the following contents: [d a:0 r:2]. This shows that the inode is of type “directory” (d, or f it is a regular file), it points to one data block (a:0, where a is for address; 0 means data block 0), and its reference count is 2 (r:2). After this, you perform certain operations on the empty root file system, but the contents are sometimes lost. Your job is to reconstruct the contents of the VSFS image.

Question 31: Assume you call `mkdir(“/n”)` on the empty root file system. The inode bitmap is missing; what should it look like:
inode bitmap ????????

```
Inodes [d a:0 r:3] [d a:1 r:2] [] [] [] [] []
data bitmap 11000000
Data [(.,0) (.,0) (n,1)] [(.,1) (.,0)] [] [] [] [] []
```

- a) 10000000
- b) 11000000
- c) 11100000
- d) 10100000
- e) None of the above

Question 32: Assume you had instead called `creat(“/z”)` on an empty file system. Unfortunately, in this case, the data block for the root directory has gone bad. What should be in there?

```
inode bitmap 11000000
Inodes [d a:0 r:2] [f a:- r:1] [] [] [] [] []
data bitmap 10000000
```

```
Data [CORRUPT!] [] [] [] [] []
```

- a) [(.,0) (.,0)]
- b) [(.,0) (.,0) (z,1)]
- c) [(.,0) (.,0) (z,2)]
- d) [(.,0) (.,0) (/z,1)]
- e) [(.,0) (.,0) (/z,2)]

This file system represents the on-disk state after two operations:

```
inode bitmap 11100000
inodes [d a:0 r:4] [d a:1 r:2] [d a:2 r:2] [] [] [] []
data bitmap 11100000
data [(.,0) (.,0) (d,1) (w,2)] [????] [(.,2) (.,0)] [] [] [] []
```

Question 33: What are the contents of the missing data block (Block 1)?

- a) (.,0) (.,0)
- b) (.,1) (.,0)
- c) (.,2) (.,0)
- d) foofoofoo
- e) None of the above

Question 34: Which two operations were run upon the empty file system to result in this state?

- a) `creat(“/d”); creat(“/w”);`
- b) `creat(“/d”); link(“/d”, “/w”);`
- c) `creat(“/w”); unlink(“/d”);`
- d) `mkdir(“/d”); mkdir(“/w”);`
- e) None of the above

Let's examine one particular corrupt file system image from VSFS for the next two questions:

inode bitmap 11000000

Inodes [d a:0 r:2] [f a:1 r:2] [] [] [] [] []

data bitmap ????????

Data [(.,0) (...0) (c,1) (m,1)] [foofoofoo] [] [] [] [] []

Question 35: Which ONE of the following is not true about the above file system state?

- a) There is a proper root directory
- b) The file "/c" exists
- c) The file "/m" exists
- d) If you read the first block of "/c", you get "foofoofoo"
- e) If you unlink "/c", you can no longer read "foofoofoo"

Question 36: Which ONE of the following is true about the above file system state? a) File "/m" is a hard link to "/c"

- b) File "/c" is a hard link to "/m"
- c) Both "/c" and "/m" are links to the same file
- d) The root directory has many data blocks in it
- e) The file "/m" has many data blocks in it

inode bitmap 11111100

Inodes [d a:0 r:3] [f a:1 r:1] [f a:2 r:1] [d a:3 r:2] [f a:- r:1] [f a:4 r:1] [] []

data bitmap 11111000

Data [(.,0) (...0) (d,1) (g,2) (y,3) (j,5)] [d] [v] [(.,3) (...0) (n,4)] [/u] [] [] []

Question 37: In this final file system state, which regular files exist?

- a) Only /d, /g, /j
- b) Only /d, /g, /j, /n
- c) Only /d, /g, /j, /y/n
- d) Only /d, /g, /j, /y e) None of the above

Question 38: In the above file system format, what is the largest number of regular-file inodes that can be allocated?

- a) 1
- b) 6
- c) 7
- d) 8
- e) As many as needed

Question 39: Bitmaps are useful as allocation structures because of all the reasons below EXCEPT:

- a) They are compact
- b) They are human readable
- c) They allow for quick lookup of free space
- d) They allow for lookup to readily find consecutive free blocks e) Updates to them do not add any disk traffic

Question 40: VSFS has all the following features EXCEPT:

- a) Regular files
- b) Directories
- c) Hard links
- d) Simple allocation structures e) Fast crash consistency

You are now given some tasks about a journaling file system. The first task is to figure out what blocks would end up in a journal transaction, given some base knowledge of the system. We assume for these first questions that the system under inspection uses data journaling mode, in which all blocks (metadata and data) are first journaled before being updated in place. Assume the standard structures of a file system here: an inode bitmap, a data bitmap, a table of inodes, and data blocks.

Important: For the questions below, ignore any additional journal metadata that would be written (i.e., a transaction start and end block), and assume no reads to disk need to take place to complete the given action (i.e., relevant structures are cached in memory).

Question 41: Assume that a process appends a data block to an existing (small) file. How many blocks are written to the journal as part of this update?

- a) 1
- b) 2
- c) 3
- d) 4
- e) None of the above

Question 42: Now assume that a process reads a block from a file. Reading, in this file system, updates the “last accessed time” field in the inode. How many blocks are written to the journal as part of this read?

- a) 1
- b) 2
- c) 3
- d) 4
- e) None of the above

Question 43: Now a process creates a 0-byte file in the root directory (which does not have many entries in it, so there is room for another entry in an existing directory data block). How many blocks are written to the journal as part of this file creation?

- a) 1
- b) 2
- c) 3
- d) 4
- e) None of the above

Question 44: Finally, a process deletes a 1-byte file from the root directory (leaving the root directory empty). Assuming the root directory only uses a single data block for its data, how many blocks are written to the journal as part of this file creation?

- a) 1
- b) 2
- c) 3
- d) 4
- e) None of the above

For the next questions, we will now assume ordered mode journaling, which only writes metadata to the journal (user file data is written only once as a result). We otherwise make the same assumptions as above, and answer the same questions about ordered mode journaling.

Question 45: Assume that a process has appends a data block to an existing (small) file. How many blocks are written to the journal as part of this update?

- a) 1
- b) 2
- c) 3
- d) 4
- e) None of the above

Question 46: Now assume that a process reads a block from a file. Reading, in this file system, updates the “last accessed time” field in the inode. How many blocks are written to the journal as part of this read?

- a) 1
- b) 2
- c) 3
- d) 4
- e) None of the above

Question 47: Now a process creates a 0-byte file in the root directory (which does not have many entries in it, so there is room for another entry in an existing directory data block). How many blocks are written to the journal as part of this file creation?

- a) 1
- b) 2
- c) 3
- d) 4
- e) None of the above

Question 48: Finally, a process deletes a 1-byte file from the root directory (leaving the root directory empty). Assuming the root directory only uses a single data block for its data, how many blocks are written to the journal as part of this file creation?

- a) 1
- b) 2
- c) 3
- d) 4
- e) None of the above

Question 49: Which of the following statements is NOT true about journaling file systems?

- a) Journaling adds a new on-disk structure to the file system
- b) Journaling is the same as write-ahead logging (the terms are used interchangeably)
- c) Journaling generally increases the amount of write traffic to the disk
- d) Journaling always makes performance worse (than the same file system without journaling)
- e) Whether in data or ordered journaling modes, file system metadata is always first written to the journal before being updated in place.

Question 50: Which of the following best represents a final, complete, and most optimized version of the ordered (metadata only) journaling protocol?

- a) Data write, then journal metadata write, then journal commit.
- b) Data write, then journal metadata write, then journal commit, then checkpoint of metadata.
- c) Data write, then journal metadata write, then journal commit, then checkpoint of metadata, then (later) mark the transaction free in the journal superblock.
- d) Data write and journal metadata write (concurrently), then journal commit, then checkpoint of metadata, then (later) mark the transaction free in the journal superblock.
- e) None of the above

You discover a disk filled with what seems to be a log-structured file system. In its initial state on disk, the first few blocks seem to be filled with an empty file system:

```
[0] checkpoint: 3 -- -- -- -- --
[1] [.,0] [...,0] -- -- -- -- --
[2] type:dir size:1 refs:2 ptrs: 1 -- -- -- -- --
[3] chunk(imap): 2 -- -- -- -- --
```

As you examine these first few blocks, you realize their structure. The first block (Block 0) is the “checkpoint region” - it points to pieces of the inode map. Block 3 holds that first chunk of the inode map; in this case, you figure out that only inode 0 is live, and it lives in Block 2. Block 2 holds the inode of the root directory, which contains 1 block, which is located in Block 1 (the first ptr in the list of addresses held in the inode). Finally, the contents of the root directory are in Block 1: a . and .. entry, each referring to the root inode 0. Good job!

Now, for some questions. You perform a single file system operation, with the resulting on-disk state:

```
[0] checkpoint: 7 -- -- -- -- -- 8 -- -- -- -- --
[1] [.,0] [...,0] -- -- -- -- --
[2] type:dir size:1 refs:2 ptrs: 1 -- -- -- -- --
[3] chunk(imap): 2 -- -- -- -- --
[4] [.,0] [...,0] [ku3,122] -- -- -- -- --
[5] type:dir size:1 refs:2 ptrs: 4 -- -- -- -- --
[6] type:reg size:0 refs:1 ptrs: -- -- -- -- --
[7] chunk(imap): 5 -- -- -- -- --
[8] chunk(imap): -- -- -- -- -- 6 -- -- -- -- --
```

Question 51: What file operation was performed?

- a) mkdir(“ku3”);
- b) creat(“ku3”);
- c) rmdir(“/”);
- d) link(“.”, “ku3”);
- e) None of the above

Question 52: In the above file system state, which of the blocks are live?

- a) 0 through 8
- b) 4 through 8
- c) 0 through 3
- d) 0 and 4 through 8
- e) None

Question 53: In the above file system state, which inodes are allocated?

- a) 7 and 8
- b) 1, 2, and 4
- c) 0 and 122
- d) 5 and 6
- e) None of the above


```
[11] chunk(imap):-----10-----
```

e) A file unlink

e) 7 through 11

e) 4 or more

e) 4 or more

```
[ 0] checkpoint: 20 -- - - - - - - - - - -  
[ 1] [.,0] [...,0] -- - - - - - - -  
[ 2] type:dir size:1 refs:2 ptrs: 1 -- - - - - - - -  
[ 3] chunk(imap): 2 -- - - - - - - - - - -  
[ 4] [.,0] [...,0] [ym6,1] -- - - - - - - -  
[ 5] type:dir size:1 refs:2 ptrs: 4 -- - - - - - - -  
[ 6] type:reg size:0 refs:1 ptrs: -- - - - - - - -  
[ 7] chunk(imap): 5 6 -- - - - - - - - - - -  
[ 8] [.,0] [...,0] [ym6,1] [hs2,2] -- - - - - -  
[ 9] type:dir size:1 refs:2 ptrs: 8 -- - - - - - - -  
[10] live type:reg size:0 refs:1 ptrs: -- - - - - - - -  
[11] chunk(imap): 9 6 10 -- - - - - - - - - - -  
[12] [.,0] [...,0] [ym6,1] [hs2,2] [cu3,1] -- - - -  
[13] type:dir size:1 refs:2 ptrs: 12 -- - - - - - - -  
[14] type:reg size:0 refs:2 ptrs: -- - - - - - - -  
[15] chunk(imap): 13 14 10 -- - - - - - - - - - -  
[16] f0f0f0f0f0f0f0f0f0f0f0f0f0f0f0f0f0  
[17] llllllllllllllllllllllllllllllllll  
[18] h2h2h2h2h2h2h2h2h2h2h2h2h2h2h2h2h2h2  
[19] type:reg size:8 refs:2 ptrs: -- - - - - - - - 16 17 18  
[20] chunk(imap): 13 19 10 -- - - - - - - - - - -
```

- a) A directory creation
- b) A file creation
- c) A multi-block write
- d) A file link
- e) None of the above

- a) A directory creation
- b) A file creation
- c) A multi-block write
- d) A file link
- e) None of the above

a) f0f0f0f0f0f0f0f0f0f0f0f0f0f0f0f0
b) llllllllllllllllllllllllllllllll
c) h2h2h2h2h2h2h2h2h2h2h2h2h2h2h2
d) All of the above
e) None of the above

a) aaaa
b) bbbb
c) cccc
d) dddd
e) None of the above

a) aaaa
b) bbbb
c) cccc
d) dddd
e) None of the above

Question 63: Assume a page-mapped FTL. The user issues a read which returns “ffff”. Which of the following could NOT accurately represent the contents of the entire FTL?

- a) 1000:1, 1001:3, 1002:5, 1003:7
- b) 1000:1, 1001:3, 1002:5, 1003:7, 1004:0
- c) 1000:1, 1001:3, 1002:5, 1003:7, 1004:0, 1005:2
- d) 1000:1, 1001:3, 1002:5, 1003:7, 1004:0, 1005:2, 1006:4
- e) None of the above

Question 64: Assume a page-mapped FTL. The user issues a read which returns “aaaa”. Which of the following could NOT accurately represent the contents of the entire FTL?

- a) 1000:1, 1001:3, 1002:5, 1003:7
- b) 1000:1, 1001:3, 1002:5, 1003:7, 1004:0
- c) 1000:1, 1001:3, 1002:5, 1003:7, 1004:0, 1005:2
- d) 1000:1, 1001:3, 1002:5, 1003:7, 1004:0, 1005:2, 1006:4
- e) None of the above

Assume the same flash contents, now showing Blocks 0, 1, and part of Block 2: [0] aaaa

[1] bbbb

[2] cccc

[3] dddd

[4] eeee

[5] ffff

[6] gggg

[7] hhhh

[8] erased (not yet programmed) [9] erased (not yet programmed) [10] etc.

For the next questions, assume the following page-mapped FTL is the initial state of the system: 10:0, 11:1, 12:2, 13:3, 14:4, 15:5, 16:6, 17:7

Question 65: Assume the user then issues a write to the SSD(address=14, data=iiii). What will the contents of flash physical page 4 be just after this write takes place?

- a) dddd
- b) eeee
- c) ffff
- d) iiii
- e) None of the above

Question 66: Assuming Block 2 (pages 8...11) of the flash is used for the write above to address=14, which of the following could represent the contents of the FTL after the write completes?

- a) 10:0, 11:1, 12:2, 13:3, 14:4, 15:5, 16:6, 17:7
- b) 10:0, 11:1, 12:2, 13:3, 14:2, 15:5, 16:6, 17:7
- c) 10:0, 11:1, 12:2, 13:3, 14:8, 15:5, 16:6, 17:7
- d) 10:0, 11:1, 12:2, 13:3, 15:5, 16:6, 17:7
- e) None of the above

Question 67: After the write (to address=14, data=iiii) is complete, the SSD decides to perform garbage collection (GC). Which of the following will NOT happen as part of the GC?

- a) Block 0 is erased
- b) Block 1 is erased
- c) Data from pages 5, 6, 7 are copied elsewhere
- d) The contents of the FTL (the mappings) will change
- e) Block 2 is erased

Question 68: Assume you have a page-mapped FTL. If each entry in the FTL takes 4 bytes (assuming it is an array), how large is the FTL? Assume the SSD is only 1 MB in size, and uses 1 KB pages.

- a) 1 KB
- b) 4 KB
- c) 1 MB
- d) 4 MB
- e) None of the above

Question 69: Assume the same system as in Question 68, but now with block mappings (not page). Assume each block fits 4 pages. Assuming each entry is still 4 bytes, how large is this block-mapped FTL?

- a) 1 KB
- b) 4 KB
- c) 1 MB
- d) 4 MB
- e) None of the above

And now, a final, general question for you about SSDs.

Question 70: Which of the following is NOT true about flash-based SSDs?

- a) SSDs need less memory to perform logical to physical translation than hard drives
- b) SSDs are generally faster than hard drives for write-oriented workloads
- c) SSDs are generally more expensive than hard drives (cost per byte)
- d) SSDs are generally faster than hard drives for read-oriented workloads
- e) SSDs use fewer moving parts than hard drives

You are given a system with 64 bytes of physical memory, 4 byte pages, and 16-byte virtual address spaces. Your forensics tools dig up the following page table structure (high bit: Valid/NOT, rest is the PFN):

```
[0] 0x00000000
[1] 0x800000?? <- missing!
[2] 0x00000000
[3] 0x00000000
```

Question 71: A trace you have accesses virtual address 0x7, which translates to 0x33. What two hex digits are missing from page table entry 1 above?

- a) 0x0a
- b) 0x0b
- c) 0x0c
- d) 0x0d
- e) None of the above

Assuming the same system (4 byte pages, 16 byte virtual address space), and the following page table:

```
[0] 0x00000000 [1] 0x80000005 [2] 0x8000000a [3] 0x00000000
```

Question 72: Given the virtual address 4, which decimal physical address does it translate to?

- a) 5
- b) 10
- c) 20
- d) 54
- e) 45

Assuming the following page table (on the same system, again):

```
[0] 0x00000000 [1] 0x8000000e [2] 0x00000000 [3] 0x80000009
```

Question 73: Which ranges of virtual address are valid? a)

- 0 ... 64
- b) 0 ... 16
- c) 4 ... 7 and 12 ... 15
- d) 0 ... 3 and 8 ... 11
- e) 80 ... 89

Question 74: You are told a given system has a 30-bit virtual address, with a 4KB page size. Assuming a 4-byte page table entry size, how big is a linear page table for a given process?

- a) 1 MB
- b) 2 MB
- c) 4 MB
- d) 8 MB
- e) None of the above

Question 75: You are next given system has a 10-bit virtual address. with a 256 byte page size. Assuming a 4-byte page table entry size, how big is a linear page table for a given process?

- a) 16 bytes
- b) 16 KB
- c) 16 MB
- d) 16 GB
- e) None of the above

Question 76: You are now given some new information about a particular system. Specifically, this system has 1 MB linear page table size (per process), and has a 1KB page size. Assuming page table entry size is 4 bytes, how many bits are in the virtual page number (VPN) on this system?

- a) 28
- b) 18
- c) 8
- d) 32
- e) None of the above

Lastly, you are given some memory dumps, which tell you something about the contents of memory. In this dump, you find that the system only has 3 pages, that pages 3, 4, and 5 are in memory. You also check a history log and find that the last 10 pages accessed were 8, 7, 4, 2, 5, 4, 7, 3, 4, 5 (in that order, with 5 being most recently accessed).

Question 77: You are then asked to determine which replacement policy was used. Is it:

- a) FIFO
- b) LRU
- c) MRU
- d) LFU
- e) Cannot tell from the given information

With the same situation as above (3-page system, with 3, 4, 5 in memory, and the last 10 pages accessed are 8, 7, 4, 2, 5, 4, 7, 3, 4, 5), you are asked the following few questions.

Question 78: Assuming the replacement policy was FIFO, how many misses were encountered while those last 10 pages were accessed? (assume the memory was empty to begin)

- a) 7
- b) 8
- c) 9
- d) 10
- e) None of the above

Question 79: Assuming the replacement policy was LRU, how many misses were encountered while those last 10 pages were accessed? (assume the memory was empty to begin)

- a) 7
- b) 8
- c) 9
- d) 10
- e) None of the above

Question 80: Assuming the replacement policy was OPT (the optimal replacement policy), how many misses were encountered while those last 10 pages were accessed? (assume the memory was empty to begin)

- a) 7
- b) 8
- c) 9
- d) 10
- e) None of the above